

STA457 Project

Xing Yu Wang

2025-03-29

```
# install.packages("forecast")
# install.packages("astsa")

library(dplyr)
library(tidyverse)
library(readr)
library(lubridate)
library(forecast)

## Warning: package 'forecast' was built under R version 4.3.3
library(astsa)

## Warning: package 'astsa' was built under R version 4.3.3
library(tseries)

## Warning: package 'tseries' was built under R version 4.3.3
library(mgcv)
library(vars)

## Warning: package 'vars' was built under R version 4.3.3
## Warning: package 'strucchange' was built under R version 4.3.3
## Warning: package 'zoo' was built under R version 4.3.3
## Warning: package 'sandwich' was built under R version 4.3.3
## Warning: package 'urca' was built under R version 4.3.3
## Warning: package 'lmtest' was built under R version 4.3.3
```

1. EDA

```
price = read.csv("./Daily Prices_ICCO.csv")
weather = read.csv("./Ghana_data.csv")
USD_GHS_Historical_Data = read.csv("./USD_GHS Historical Data.csv")
```

1.1 Clean Data

```
weather <- weather |> dplyr::select(DATE, TAVG)
exchangerate <- USD_GHS_Historical_Data |> dplyr::select(Date, Price)
```

```
colnames(price)[colnames(price) == 'ICCO.daily.price..US..tonne.'] <- 'Daily_Price'
colnames(weather)[colnames(weather) == 'DATE'] <- 'Date'
colnames(weather)[colnames(weather) == 'TAVG'] <- 'Avg_Temp'
colnames(exchangerate)[colnames(exchangerate) == 'Price'] <- 'exchange_rate'
```

1.2 Check duplicated values

```
price |> group_by(Date) |> filter(n() > 1) |> ungroup()
```

```
## # A tibble: 8 x 2
##   Date       Daily_Price
##   <chr>      <chr>
## 1 31/01/2024 4,798.20
## 2 31/01/2024 10,888.05
## 3 30/01/2024 4,775.17
## 4 30/01/2024 10,676.42
## 5 09/01/2024 4,171.24
## 6 09/01/2024 4,171.24
## 7 15/12/2023 4,272.15
## 8 15/12/2023 4,272.15
```

```
price <- price |> filter(!(Date == "31/01/2024" & Daily_Price == "10,888.05"))
price <- price |> filter(!(Date == "30/01/2024" & Daily_Price == "10,676.42"))
price <- distinct(price)
```

1.3 Convert to Time Series Data

1.3.1 price Dataset

```
price$Date <- as.Date(price$Date, format="%d/%m/%Y")
price$Daily_Price <- as.numeric(gsub(",", "", price$Daily_Price))
price_month <- price |> mutate(Time = floor_date(Date, "month")) |> group_by(Time) |>
  summarise(month_Price = mean(Daily_Price, na.rm = TRUE)) |> ungroup()
```

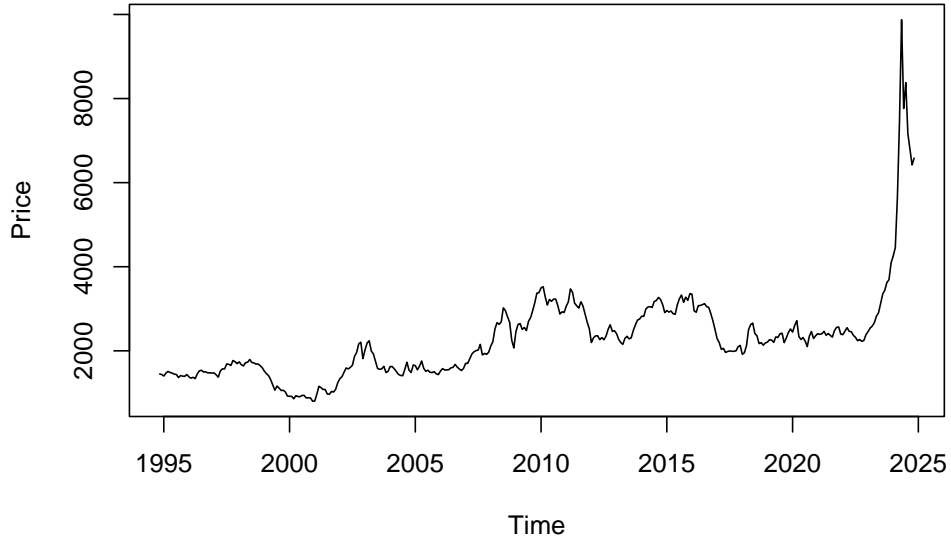
```
summary(price)
```

```
##      Date       Daily_Price
## Min.   :1994-10-03 Min.    : 774.1
## 1st Qu.:2002-05-16 1st Qu.: 1557.8
## Median :2009-12-17 Median  : 2202.0
## Mean   :2009-12-17 Mean    : 2350.1
## 3rd Qu.:2017-07-24 3rd Qu.: 2738.1
## Max.   :2025-02-27 Max.    :11984.7
```


```
price_ts <- ts(price_month$month_Price, start = c(1994, 11), end = c(2024, 11), frequency = 12)
```

```
plot(price_ts, main="Monthly Price Time Series", ylab="Price", xlab="Time")
```

Monthly Price Time Series



```
acf2(price_ts, 50)
```

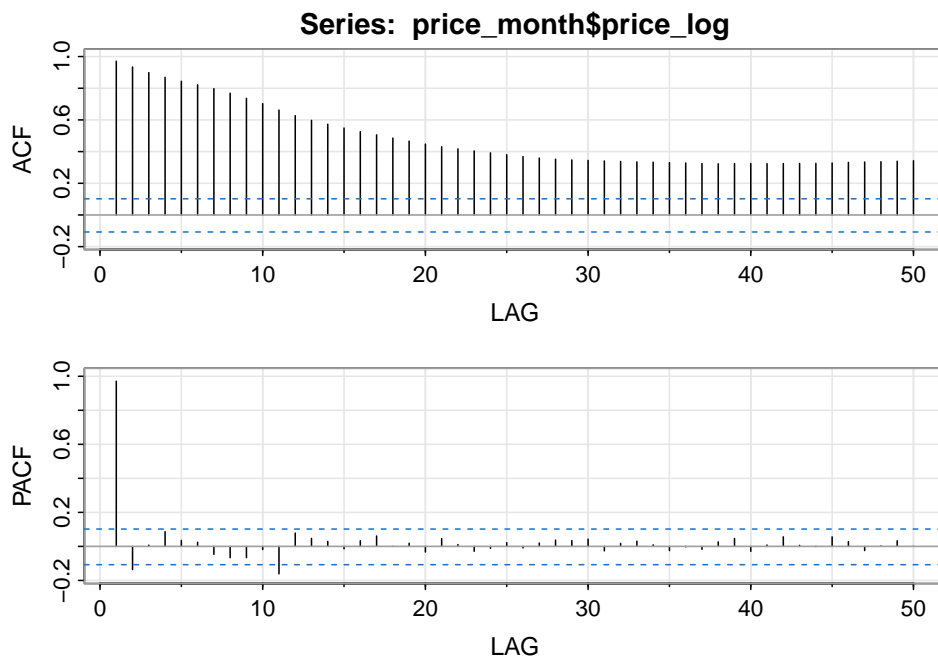


	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]	[,12]	[,13]
## ACF	0.95	0.90	0.84	0.78	0.71	0.65	0.57	0.51	0.47	0.44	0.42	0.39	0.37
## PACF	0.95	-0.04	-0.16	-0.01	-0.14	0.08	-0.24	0.18	0.23	0.00	0.01	-0.13	0.08
	[,14]	[,15]	[,16]	[,17]	[,18]	[,19]	[,20]	[,21]	[,22]	[,23]	[,24]	[,25]	
## ACF	0.35	0.33	0.31	0.30	0.28	0.27	0.26	0.25	0.25	0.24	0.23	0.23	
## PACF	-0.07	-0.07	0.09	0.05	0.10	-0.07	0.00	0.03	-0.04	0.00	-0.02	0.10	
	[,26]	[,27]	[,28]	[,29]	[,30]	[,31]	[,32]	[,33]	[,34]	[,35]	[,36]	[,37]	
## ACF	0.23	0.22	0.22	0.22	0.22	0.21	0.21	0.20	0.20	0.20	0.20	0.19	
## PACF	0.02	-0.04	0.04	0.00	0.03	-0.08	0.00	0.05	0.01	0.03	-0.02	0.04	
	[,38]	[,39]	[,40]	[,41]	[,42]	[,43]	[,44]	[,45]	[,46]	[,47]	[,48]	[,49]	
## ACF	0.19	0.19	0.19	0.18	0.18	0.19	0.18	0.18	0.18	0.18	0.19	0.19	
## PACF	-0.02	-0.02	0.01	0.00	0.06	0.00	-0.04	0.03	0.04	-0.01	0.00	0.05	

```
##      [,50]
## ACF    0.20
## PACF   0.03
ndiffs(price_ts)
```

```
## [1] 1
price_month$price_log <- log(price_month$month_Price)
adf.test(price_month$price_log)
```

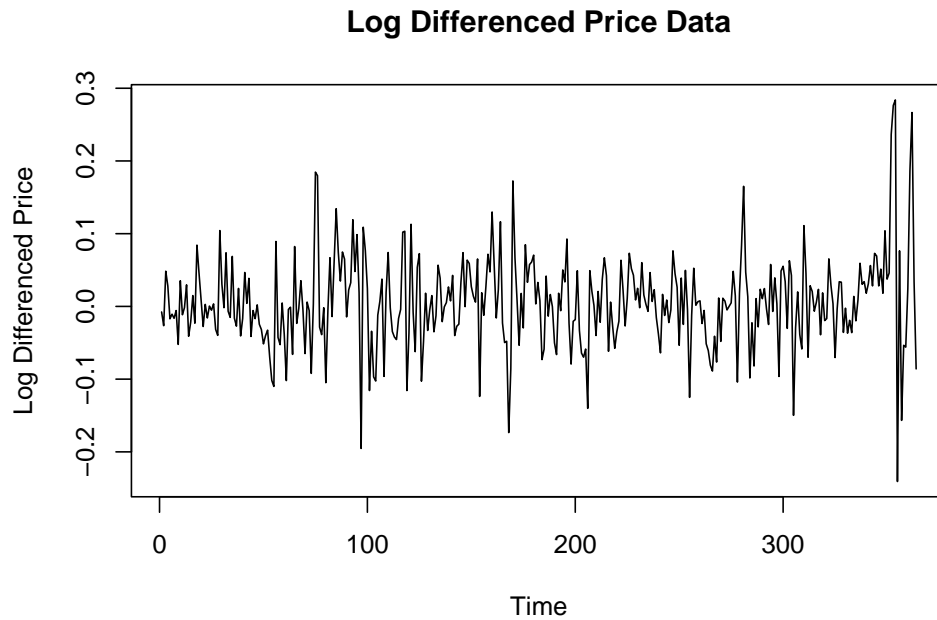
```
##
## Augmented Dickey-Fuller Test
##
## data: price_month$price_log
## Dickey-Fuller = -1.736, Lag order = 7, p-value = 0.6883
## alternative hypothesis: stationary
acf2(price_month$price_log, 50)
```



```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13]
## ACF  0.97  0.93  0.90  0.87  0.84  0.82  0.80  0.77  0.74  0.70  0.66  0.63  0.60
## PACF  0.97 -0.14  0.01  0.09  0.04  0.02 -0.05 -0.07 -0.07 -0.02 -0.16  0.08  0.05
##      [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25]
## ACF  0.57  0.55  0.53  0.51  0.49  0.47  0.45  0.43  0.42  0.40  0.39  0.38
## PACF  0.03 -0.01  0.03  0.06  0.00  0.02 -0.03  0.05  0.01 -0.03 -0.01  0.02
##      [,26] [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37]
## ACF  0.37  0.36  0.35  0.35  0.34  0.34  0.34  0.34  0.33  0.33  0.33  0.32
## PACF -0.01  0.02  0.04  0.03  0.04 -0.03  0.02  0.03  0.01 -0.02  0.00 -0.02
##      [,38] [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49]
## ACF  0.32  0.32  0.32  0.32  0.32  0.33  0.33  0.33  0.33  0.33  0.34  0.34
## PACF  0.03  0.05 -0.03  0.01  0.06  0.01  0.00  0.06  0.03 -0.02  0.00  0.03
##      [,50]
## ACF    0.34
## PACF   0.00
```

Hence, we want to difference the price data.

```
diff_log_price = diff(price_month$price_log)
ts.plot(diff_log_price, main = "Log Differenced Price Data", ylab = "Log Differenced Price")
```



```
adf.test(diff_log_price)
```

```
## Warning in adf.test(diff_log_price): p-value smaller than printed p-value
```

```
##
```

```
## Augmented Dickey-Fuller Test
```


```
##
```

```
## data: diff_log_price
```

```
## Dickey-Fuller = -6.1385, Lag order = 7, p-value = 0.01
```

```
## alternative hypothesis: stationary
```

```
acf2(diff_log_price, 50)
```



```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13]
## ACF  0.22 -0.01 0.00  0.00 -0.11 0.03 0.12  0.02 0.09  0.08  0.05  0.02  0.02
## PACF 0.22 -0.06 0.01 -0.01 -0.12 0.09 0.10 -0.03 0.11  0.02  0.04  0.03 -0.01
##      [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25]
## ACF   0.02 -0.04  0  0.00  0.02  0.04 -0.03 -0.08 -0.01  0.03  0  0.01
## PACF  0.03 -0.04  0 -0.01  0.00  0.04 -0.08 -0.07  0.04  0.01  0  0.00
##      [,26] [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37]
## ACF  -0.06 -0.09 -0.07 -0.05  0.02  0.00 -0.05 -0.04 -0.02  0.01 -0.01 -0.04
## PACF -0.10 -0.03 -0.04 -0.04  0.06 -0.04 -0.06 -0.01 -0.02  0.06 -0.01 -0.05
##      [,38] [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49]
## ACF  -0.09  0.02 -0.02 -0.06  0.02  0.09 -0.11 -0.10  0.08 -0.02 -0.06  0.03
## PACF -0.06  0.08 -0.02 -0.06  0.04  0.08 -0.14 -0.01  0.11 -0.07 -0.02  0.01
##      [,50]
## ACF   0.09
## PACF  0.07
```

1.3.2 ghana Dataset

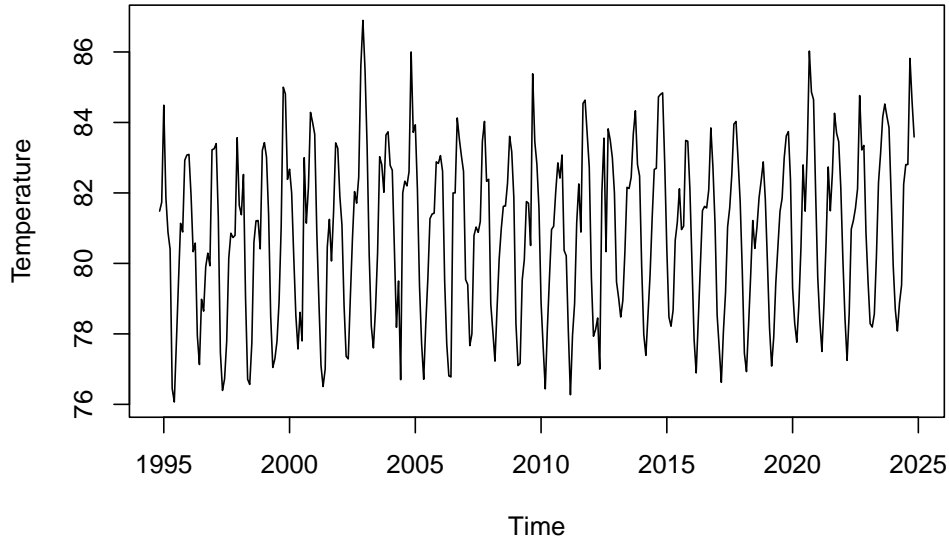
```
weather$Date <- as.Date(weather$Date)
weather$Avg_Temp <- as.numeric(gsub("", "", weather$Avg_Temp))
weather_month <- weather |> mutate(Time = floor_date(Date, "month")) |> group_by(Time) |>
  summarise(Avg_Temp = mean(Avg_Temp, na.rm = TRUE)) |> ungroup()

summary(weather_month)
```

```
##      Time      Avg_Temp
## Min.   :1990-01-01  Min.   :76.07
## 1st Qu.:1998-09-23  1st Qu.:78.90
## Median :2007-07-16  Median :81.20
## Mean   :2007-06-22  Mean    :80.97
## 3rd Qu.:2016-03-08  3rd Qu.:82.82
## Max.   :2024-11-01  Max.    :86.90
```

```
weather_ts <- ts(weather_month$Avg_Temp, start = c(1994, 11), end = c(2024, 11), frequency = 12)
ts.plot(weather_ts, main="Monthly Average Temperature Time Series", ylab="Temperature", xlab="Time")
```

Monthly Average Temperature Time Series



1.3.3 exchange Data

```
exchangerate$Date <- as.Date(exchangerate$Date)
exchangerate$exchange_rate <- as.numeric(gsub(",", "", exchangerate$exchange_rate))
rate_month <- exchangerate |> mutate(Time = floor_date(Date, "month")) |> group_by(Time) |>
  summarise(exchange_rate = mean(exchange_rate, na.rm = TRUE)) |> ungroup()
```

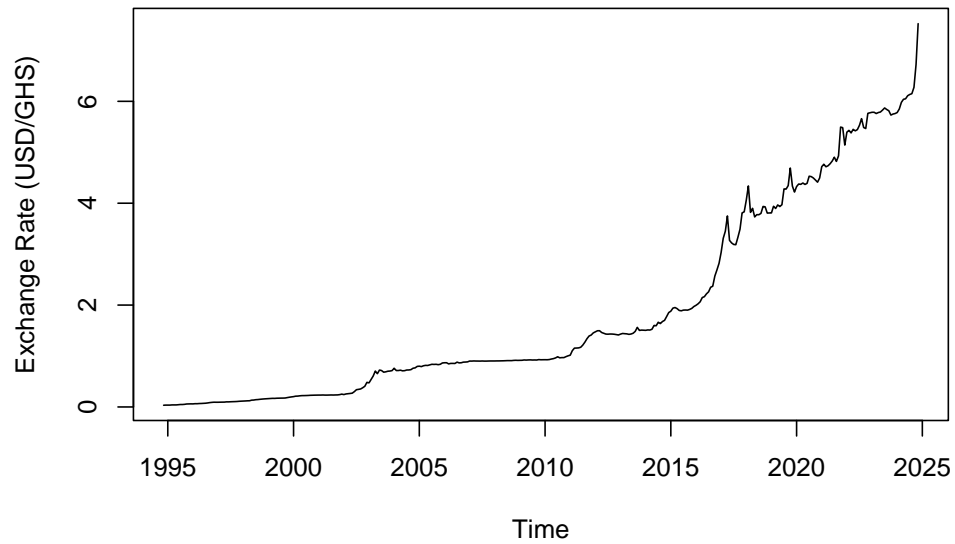
```
summary(exchangerate)
```

```
##      Date      exchange_rate
## Min.   :1992-03-01   Min.    : 0.0338
## 1st Qu.:2000-06-01   1st Qu.: 0.5400
## Median :2008-09-01   Median : 1.1595
## Mean   :2008-08-31   Mean    : 2.8314
## 3rd Qu.:2016-12-01   3rd Qu.: 4.2805
## Max.   :2025-03-01   Max.    :16.2500
```

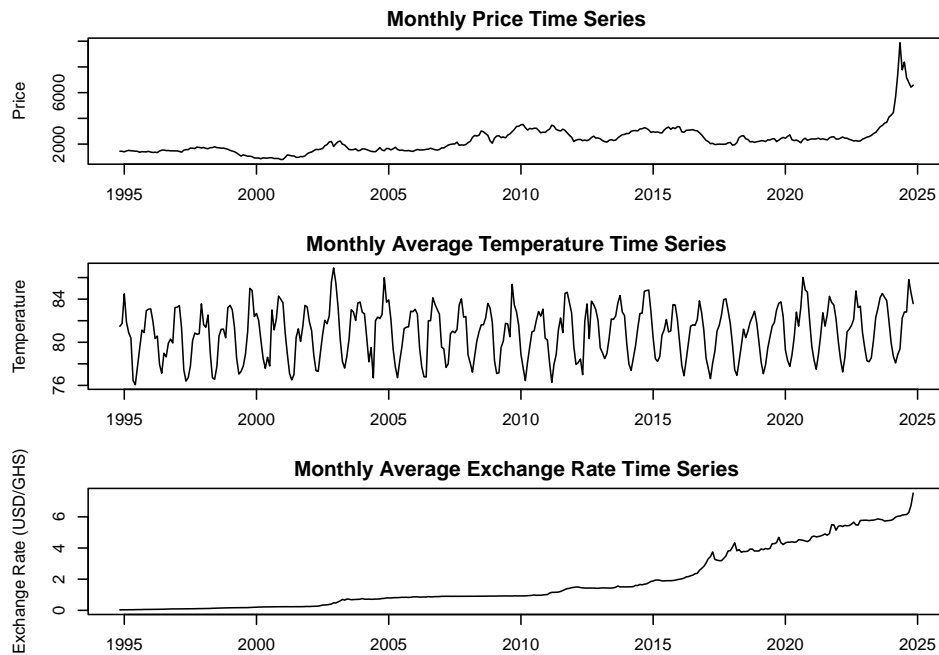
```
rate_ts <- ts(rate_month$exchange_rate, start = c(1994, 11), end = c(2024, 11), frequency = 12)
```

```
ts.plot(rate_ts, main="Monthly Average Exchange Rate Time Series", ylab="Exchange Rate (USD/GHS)", xlab="Time")
```

Monthly Average Exchange Rate Time Series

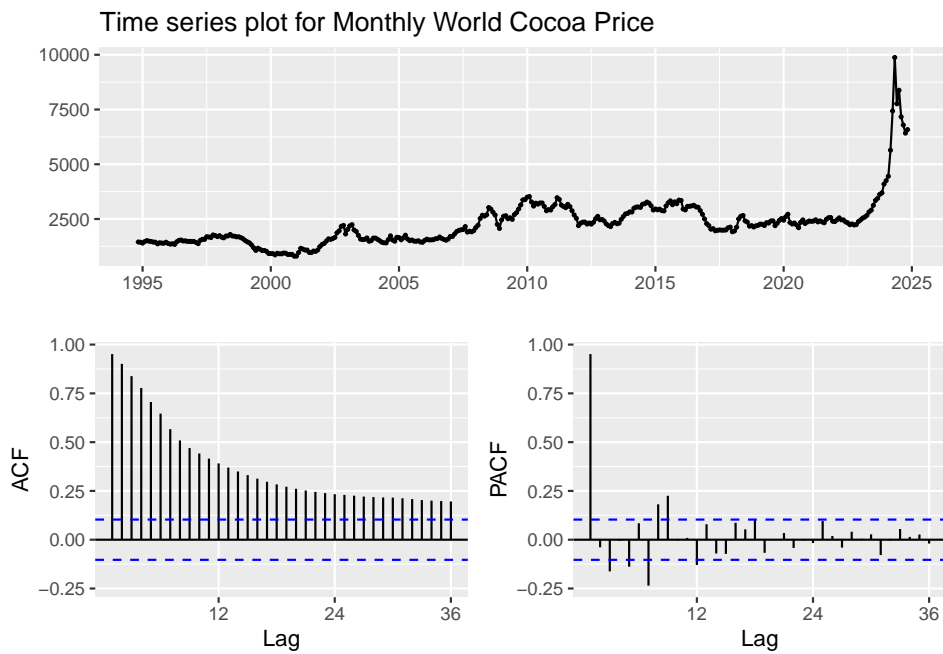


```
par(mfrow=c(3,1), mar = c(3, 4, 2, 2))
# price
plot(price_ts, main="Monthly Price Time Series", ylab="Price", xlab="Time")
# temperature
ts.plot(weather_ts, main="Monthly Average Temperature Time Series", ylab="Temperature", xlab="Time")
# exchange rate
ts.plot(rate_ts, main="Monthly Average Exchange Rate Time Series", ylab="Exchange Rate (USD/GHS)", xlab="Time")
```

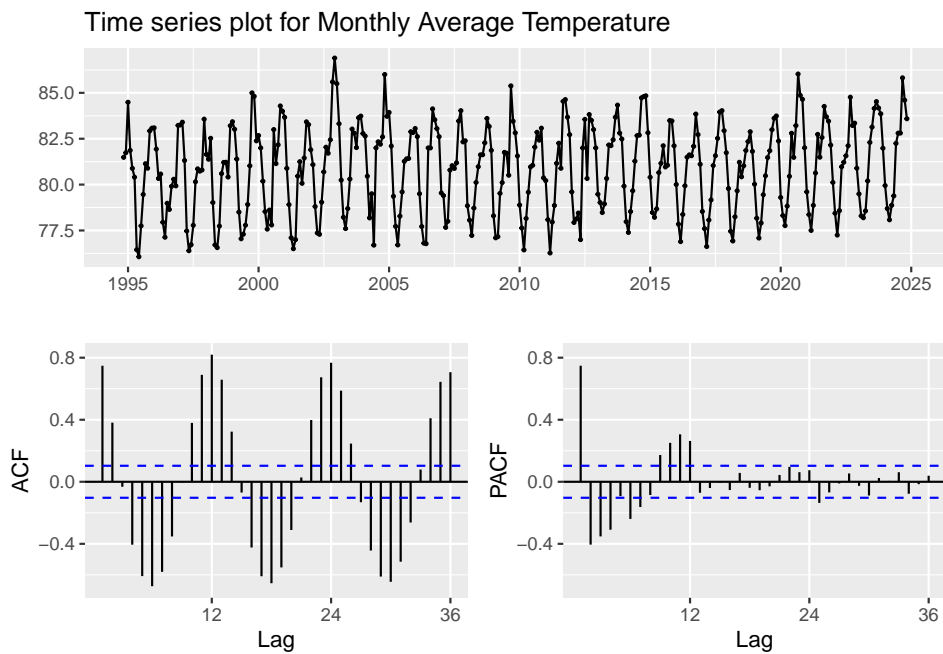


1.4 Time series plots for data


```
ggtsdisplay(price_ts, main="Time series plot for Monthly World Cocoa Price")
```

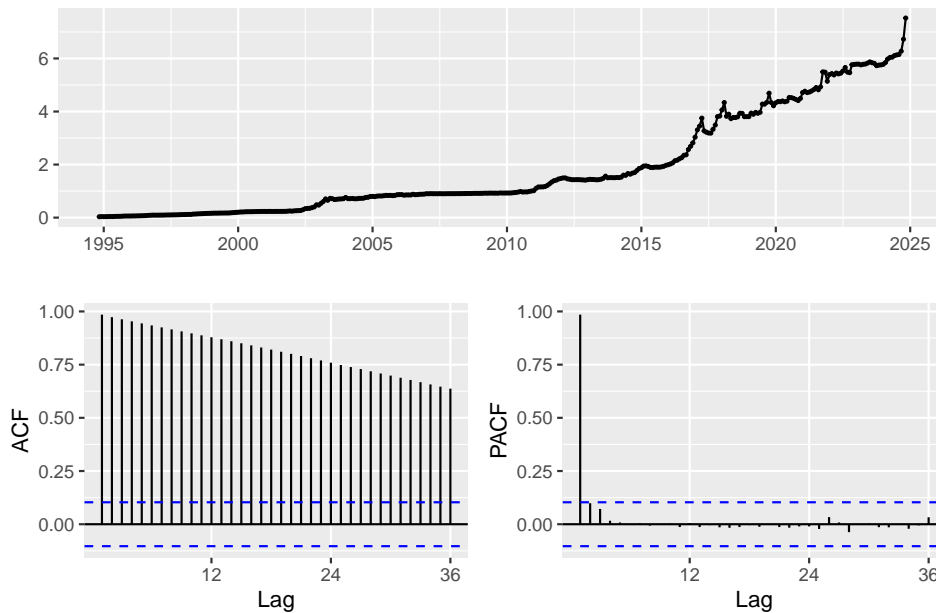


```
ggtsdisplay(weather_ts, main="Time series plot for Monthly Average Temperature")
```

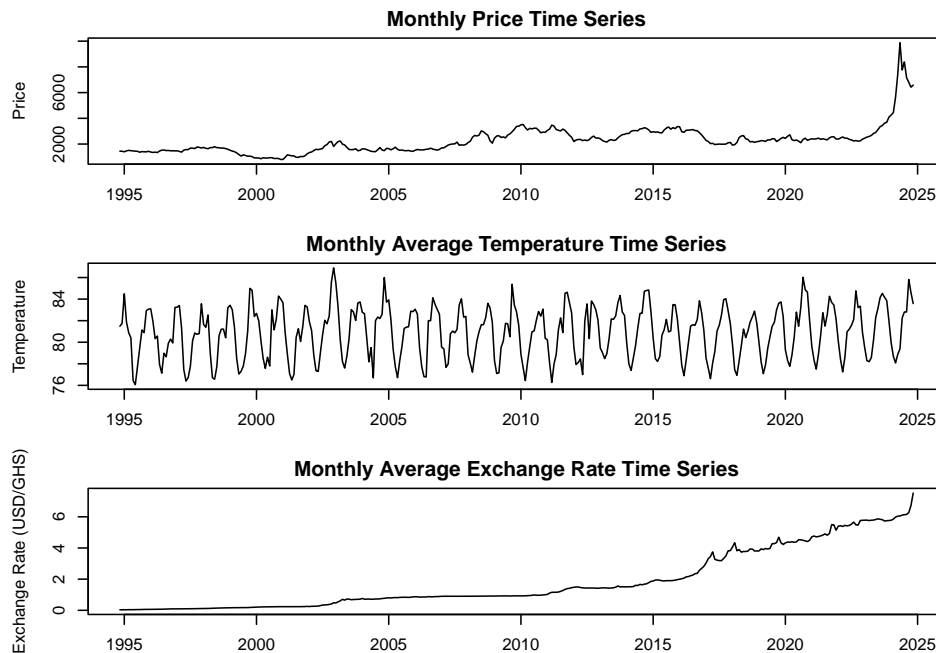


```
ggtsdisplay(rate_ts, main="Time series plot for Monthly Average Exchange Rate(USD/GHS)")
```

Time series plot for Monthly Average Exchange Rate(USD/GHS)



```
par(mfrow=c(3,1), mar = c(3, 4, 2, 2))
# price
plot(price_ts, main="Monthly Price Time Series", ylab="Price", xlab="Time")
# temperature
ts.plot(weather_ts, main="Monthly Average Temperature Time Series", ylab="Temperature", xlab="Time")
# exchange rate
ts.plot(rate_ts, main="Monthly Average Exchange Rate Time Series", ylab="Exchange Rate (USD/GHS)", xlab="Time")
```



1.5 Combine and Split data

```
data <- price_month |> left_join(weather_month, by = "Time") |> left_join(rate_month, by = "Time")
data <- data |> mutate(log_price = log(month_Price), diff_log_price =
```

```

      c(NA, diff(price_month$price_log))) |> drop_na()
data <- data |> dplyr::select(Time, Avg_Temp, exchange_rate, diff_log_price, log_price, month_Price)

data$Time <- as.Date(data$Time)

data <- data[order(data$Time), ]
cutoff <- floor(0.7 * nrow(data))
trainSet <- data[1:cutoff, ]
testSet <- data[(cutoff+1):nrow(data), ]

data_train_ts <- ts(trainSet$diff_log_price, frequency = 12)

```

1.6 Stationarity check and Decomposition

```

adf.test(data$month_Price)

##
## Augmented Dickey-Fuller Test
##
## data: data$month_Price
## Dickey-Fuller = -1.7041, Lag order = 7, p-value = 0.7017
## alternative hypothesis: stationary

adf.test(data$log_price)

##
## Augmented Dickey-Fuller Test
##
## data: data$log_price
## Dickey-Fuller = -2.3875, Lag order = 7, p-value = 0.4133
## alternative hypothesis: stationary

adf.test(data$diff_log_price)

## Warning in adf.test(data$diff_log_price): p-value smaller than printed p-value

##
## Augmented Dickey-Fuller Test
##
## data: data$diff_log_price
## Dickey-Fuller = -6.2103, Lag order = 7, p-value = 0.01
## alternative hypothesis: stationary


```

Since only the `diff_log_price` is stationary, we choose differenced monthly log price when fitting the model.

```

diff_price_ts <- ts(data$diff_log_price, frequency = 12)
autoplot(decompose(diff_price_ts, type="additive")) +
  ggtitle("Decomposition of Differenced Log Price Time Series") +
  theme_minimal()

```



2. Method

2.1 ETS Model

ETS is a purely univariate model and cannot directly handle external regressors.

2.1.1 Fit Model

```
ets_model <- ets(data_train_ts)
ets_zmodel <- ets(data_train_ts, model = "ZZZ") # Automatically selects best model
summary(ets_model)
```

```
## ETS(A,N,N)
##
## Call:
## ets(y = data_train_ts)
##
## Smoothing parameters:
##   alpha = 1e-04
##
## Initial states:
##   l = 0.0029
##
## sigma: 0.0569
##
##      AIC      AICc      BIC
## -48.96308 -48.86552 -38.39869
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 1.567182e-05 0.05666171 0.04285329 109.1957 114.6766 0.694269
##              ACF1
```

```
## Training set 0.1682833
```

```
summary(ets_zmodel)
```

```
## ETS(A,N,N)
```

```
##
```

```
## Call:
```

```
## ets(y = data_train_ts, model = "ZZZ")
```

```
##
```

```
## Smoothing parameters:
```

```
## alpha = 1e-04
```

```
##
```

```
## Initial states:
```

```
## l = 0.0029
```

```
##
```

```
## sigma: 0.0569
```

```
##
```

```
## AIC AICc BIC
```

```
## -48.96308 -48.86552 -38.39869
```

```
##
```

```
## Training set error measures:
```

```
## ME RMSE MAE MPE MAPE MASE
```

```
## Training set 1.567182e-05 0.05666171 0.04285329 109.1957 114.6766 0.694269
```

```
## ACF1
```

```
## Training set 0.1682833
```

ETS(A,N,N) is the best model.

2.1.2 Forecasting and Plotting

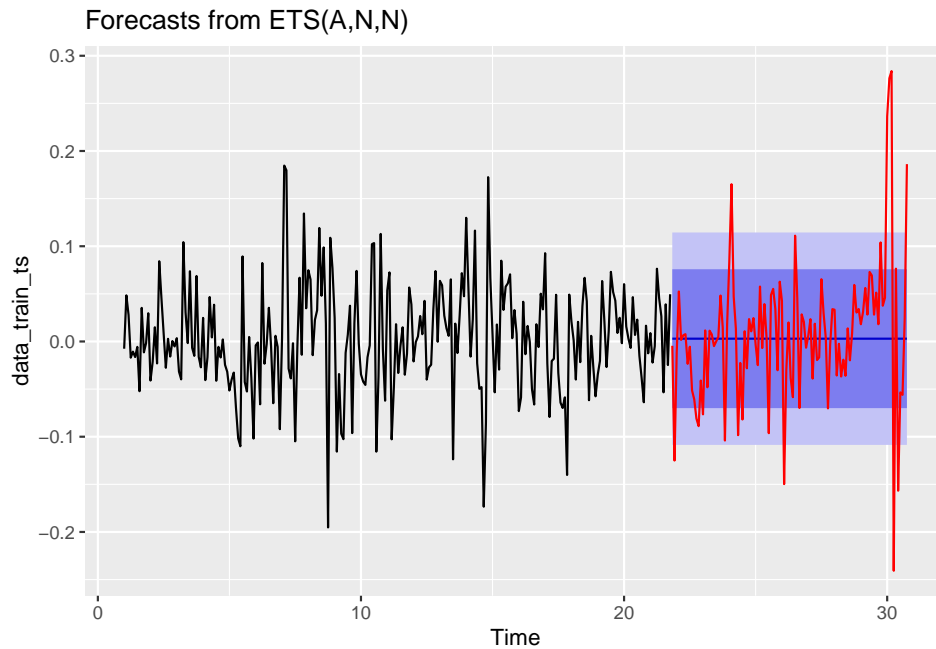
```
# Plot using log differenced price
```

```
data_test_ts <- ts(testSet$diff_log_price, start = end(data_train_ts) + c(0,1),  
frequency = 12)
```

```
h <- nrow(testSet)
```

```
forecast_ets <- forecast(ets_model, h = h)
```

```
autoplot(forecast_ets) + autolayer(data_test_ts, series = "Actual", color = "red")
```



The red line is the observed actual values. The forecasted values are the central blue line within the blue shaded prediction intervals.

```
last_log_price <- tail(trainSet$log_price , 1)
forecasted_log_price <- cumsum(forecast_ets$mean) + last_log_price


# Convert back to actual price
forecasted_price <- exp(forecasted_log_price)

actual_price <- exp(testSet$log_price)

data_test_ts <- ts(testSet$diff_log_price, start = end(data_train_ts) + c(0,1),
                  frequency = 12)

forecast_ets_ts <- ts(forecasted_price, start = start(data_test_ts), frequency = 12)
actual_ets_ts <- ts(actual_price, start = start(data_test_ts), frequency = 12)

# Plot using actual price
autoplot(forecast_ets_ts, series = "Predicted") +
  autolayer(actual_ets_ts, series = "Actual", color = "blue") +
  ggtitle("Forecast vs Actual Prices") +
  ylab("Price") +
  xlab("Time") +
  theme_minimal()
```



2.2 ARIMAX Model

Recall that in Section 1.3.1, we have tested the `acf` and `adf.test`, and determined that we would be using the differenced price data. To fit the trainset, we evaluate `p` and `q` for ARIMA model.

```
adf.test(trainSet$diff_log_price)
```

```
## Warning in adf.test(trainSet$diff_log_price): p-value smaller than printed
```

```
## p-value
```

```
##
```

```
## Augmented Dickey-Fuller Test
```


```
##
```

```
## data: trainSet$diff_log_price
```

```
## Dickey-Fuller = -5.015, Lag order = 6, p-value = 0.01
```

```
## alternative hypothesis: stationary
```

```
acf2(trainSet$diff_log_price)
```



```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13]
## ACF  0.17 -0.09 0.04 -0.01 -0.10 0.09 0.15 -0.08 -0.06 -0.01 0.04 0.05 -0.05
## PACF 0.17 -0.12 0.08 -0.04 -0.09 0.13 0.09 -0.10 -0.01 -0.03 0.08 0.04 -0.11
##      [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25]
## ACF  -0.01 0.03 0.03 -0.01 0.06 0.05 -0.18 -0.03 0.08 -0.05 -0.11 0.04
## PACF  0.02 0.05 0.05 -0.02 0.02 0.05 -0.17 0.04 0.02 -0.07 -0.07 0.02
##      [,26]
## ACF  -0.09
## PACF  -0.11
```

2.2.1 Fit ARIMAX Model

```
xreg_matrix <- cbind(trainSet$Avg_Temp, trainSet$exchange_rate)

# Assign column names to xreg
colnames(xreg_matrix) <- c("Avg_Temp", "exchange_rate")

# Fit the SARIMA model with the named xreg
arimax_model <- Arima(trainSet$diff_log_price, order=c(1,0,1), xreg = xreg_matrix)

# Summary of the model
summary(arimax_model)
```

```
## Series: trainSet$diff_log_price
## Regression with ARIMA(1,0,1) errors
##
## Coefficients:
##      ar1      ma1  intercept  Avg_Temp  exchange_rate
##    -0.3267  0.5331   -0.0325   0.0004     0.0035
## s.e.   0.1996  0.1763    0.1384   0.0017     0.0044
##
## sigma^2 = 0.003118: log likelihood = 369.07
## AIC=-726.14   AICc=-725.79   BIC=-705.01
```




```
##
## Training set error measures:
##           ME           RMSE          MAE          MPE          MAPE          MASE
## Training set 1.154543e-05 0.05527999 0.04194407 130.5909 158.5406 0.7578074
##           ACF1
## Training set -0.009038859

AIC(arimax_model)

## [1] -726.1383

checkresiduals(arimax_model)
```



```
##
## Ljung-Box test
##
## data: Residuals from Regression with ARIMA(1,0,1) errors
## Q* = 13.437, df = 8, p-value = 0.09766
##
## Model df: 2. Total lags used: 10
```

Fail to reject H_0 , hence residuals of this plot do not show significant autocorrelation.

2.2.2 Forecasting With ARIMAX Model

```
# Create future xreg from test set
forecast_arimax_xreg <- cbind(testSet$Avg_Temp, testSet$exchange_rate)

# Ensure column names match the original xreg
colnames(forecast_arimax_xreg) <- c("Avg_Temp", "exchange_rate")

# Forecast using the ARIMAX model with xreg
forecast_arimax <- forecast(arimax_model, xreg = forecast_arimax_xreg)

# Assuming your model was trained on log_price, use the last log_price from the training set
```

```

last_log_price <- tail(trainSet$log_price, 1)

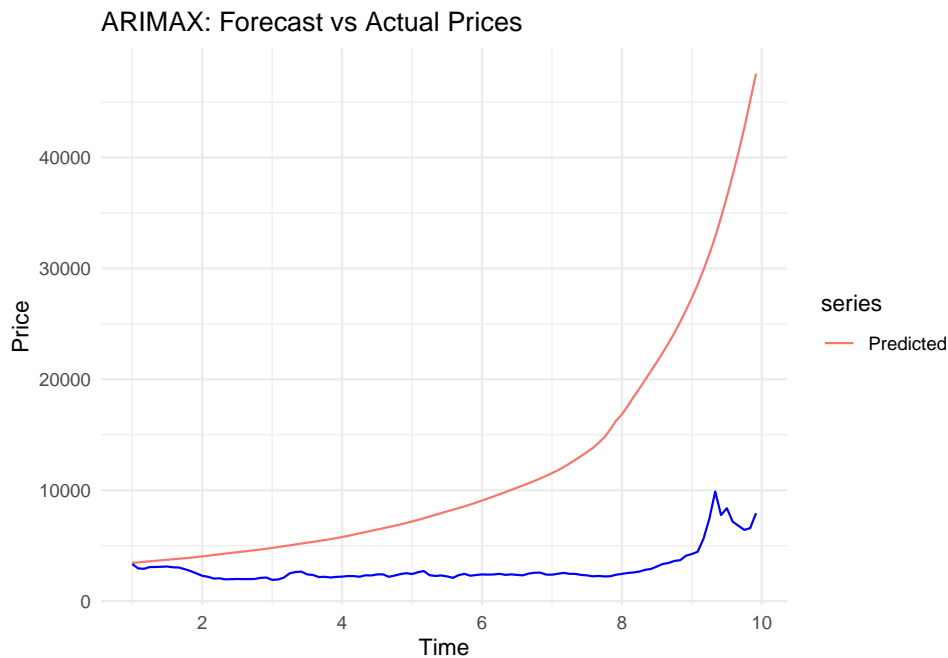
# Convert the forecasted log prices back to actual prices
forecasted_price_arimax <- exp(cumsum(forecast_arimax$mean) + last_log_price)

# Get actual prices from the test set
actual_price_arimax <- exp(cumsum(testSet$diff_log_price) + last_log_price)

# Create time series objects with the correct start point and frequency
forecast_arimax_ts <- ts(forecasted_price_arimax, start = start(testSet$Time), frequency = 12)
actual_arimax_ts <- ts(actual_price_arimax, start = start(testSet$Time), frequency = 12)

# Plot forecast vs actual prices
autoplot(forecast_arimax_ts, series = "Predicted") +
  autolayer(actual_arimax_ts, series = "Actual", color = "blue") +
  ggtitle("ARIMAX: Forecast vs Actual Prices") +
  ylab("Price") +
  xlab("Time") +
  theme_minimal()

```



2.5 GAM Model

2.5.1 Fit Model

```

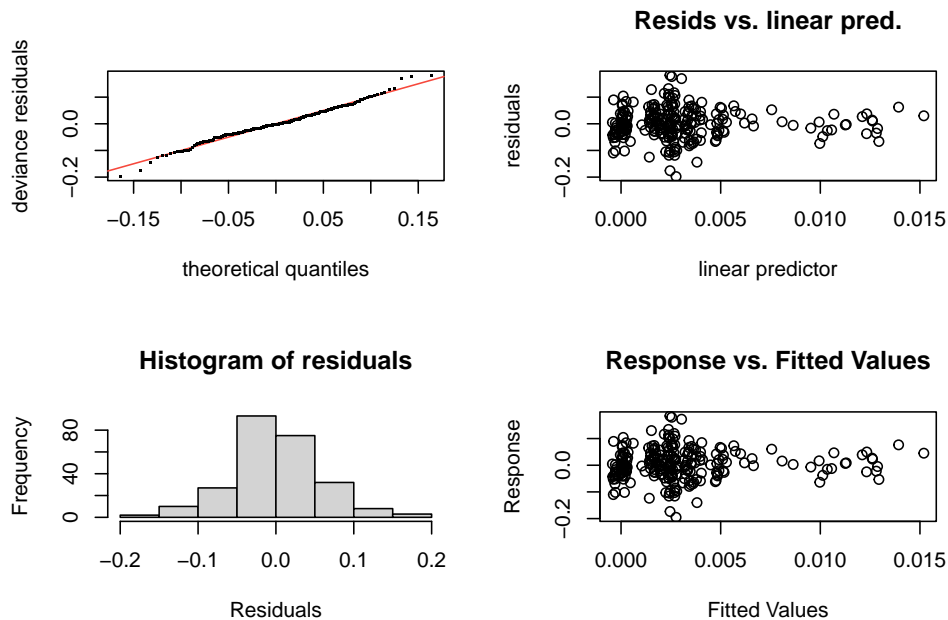
# Uses simple smoothing splines for variables
trainSet$Time_num <- as.numeric(trainSet$Time)

gam_basic <- gam(diff_log_price ~ s(Time_num) + s(Avg_Temp) + s(exchange_rate),
  data = trainSet, method = "REML")

gam.check(gam_basic)

```

2.5.1.1 Basic Model



```
##
## Method: REML   Optimizer: outer newton
## full convergence after 10 iterations.
## Gradient range [-0.0001493842,0.0002242223]
## (score -345.2659 & scale 0.003250926).
## Hessian positive definite, eigenvalue range [1.158838e-05,122.9998].
## Model rank = 28 / 28
##
## Basis dimension (k) checking results. Low p-value (k-index<1) may
## indicate that k is too low, especially if edf is close to k'.
##
##          k'   edf k-index p-value
## s(Time_num)   9.00 1.00   0.84 <2e-16 ***
## s(Avg_Temp)   9.00 1.00   1.06   0.81
## s(exchange_rate) 9.00 1.14   0.90   0.02 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Fit GAM model with smooth trend, random effects, and seasonality
trainSet$dateInt = as.integer(trainSet$Time)
trainSet$monthFac <- factor(format(trainSet$Time, "%m"))
gam_complex <- gam(diff_log_price ~ s(dateInt, k = 100) + s(monthFac, bs = "re") +
  sinpi(dateInt / 182.625) + cospi(dateInt / 182.625) +
  sinpi(dateInt / 91.3125) + cospi(dateInt / 91.3125) +
  s(Avg_Temp) + s(exchange_rate), data = trainSet,
  method = "REML", optimizer = "efs")
```

```
summary(gam_basic)
```

2.5.1.2 Complex Model

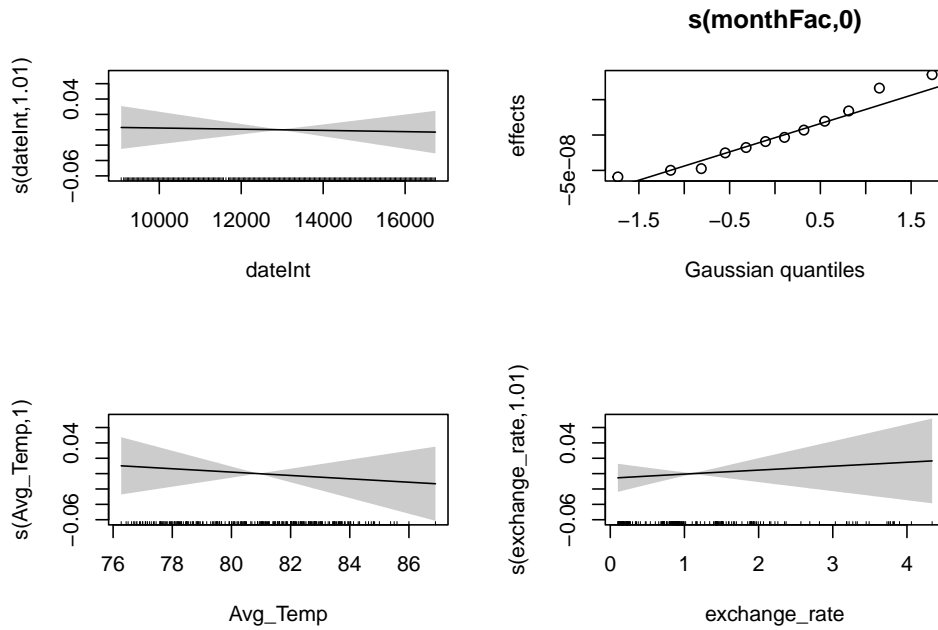
```
##
```

```

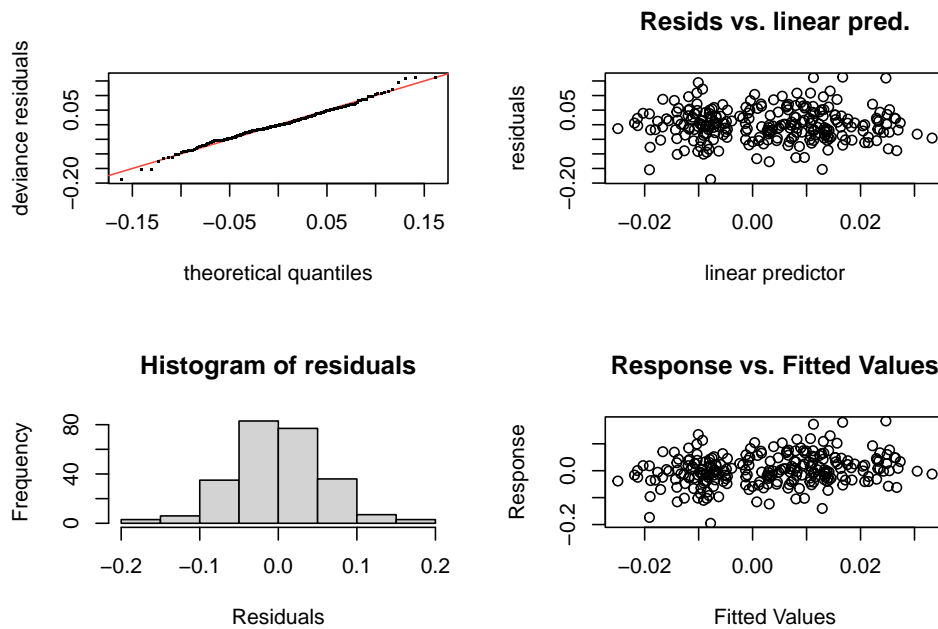
## Family: gaussian
## Link function: identity
##
## Formula:
## diff_log_price ~ s(Time_num) + s(Avg_Temp) + s(exchange_rate)
##
## Parametric coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.002929   0.003606   0.812   0.417
##
## Approximate significance of smooth terms:
##               edf Ref.df    F p-value
## s(Time_num)      1.000  1.000 0.074   0.786
## s(Avg_Temp)       1.001  1.001 0.000   0.996
## s(exchange_rate) 1.142  1.276 0.266   0.695
##
## R-sq.(adj) = -0.00863   Deviance explained = 0.41%
## -REML = -345.27   Scale est. = 0.0032509   n = 250
summary(gam_complex)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## diff_log_price ~ s(dateInt, k = 100) + s(monthFac, bs = "re") +
##   sinpi(dateInt/182.625) + cospi(dateInt/182.625) + sinpi(dateInt/91.3125) +
##   cospi(dateInt/91.3125) + s(Avg_Temp) + s(exchange_rate)
##
## Parametric coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.0029355   0.0035494   0.827   0.4090
## sinpi(dateInt/182.625) 0.0081148   0.0085424   0.950   0.3431
## cospi(dateInt/182.625) 0.0094069   0.0106187   0.886   0.3766
## sinpi(dateInt/91.3125) -0.0000812   0.0050924  -0.016   0.9873
## cospi(dateInt/91.3125) 0.0144616   0.0060586   2.387   0.0178 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##               edf Ref.df    F p-value
## s(dateInt)      1.006e+00  1.013 0.048   0.842
## s(monthFac)      4.392e-05 11.000 0.000   0.809
## s(Avg_Temp)      1.000e+00  1.000 0.297   0.586
## s(exchange_rate) 1.013e+00  1.027 0.351   0.561
##
## R-sq.(adj) = 0.0234   Deviance explained = 5.09%
## -REML = -334.76   Scale est. = 0.0031478   n = 250
plot(gam_complex, pages = 1, shade = TRUE)

```



```
gam.check(gam_complex)
```



```
##
## Method: REML   Optimizer: efs
## $iter
## [1] 31
##
## $score.hist
## [1] -332.7868 -333.5308 -334.0373 -334.3219 -334.4896 -334.5939 -334.6572
## [8] -334.6941 -334.7155 -334.7283 -334.7364 -334.7418 -334.7455 -334.7481
## [15] -334.7501 -334.7515 -334.7526 -334.7535 -334.7542 -334.7548 -334.7553
## [22] -334.7556 -334.7559 -334.7562 -334.7564 -334.7566 -334.7568 -334.7569
## [29] -334.7570 -334.7571 -334.7572
##
```

```

## $conv
## [1] "full convergence"
##
## Model rank = 134 / 134
##
## Basis dimension (k) checking results. Low p-value (k-index<1) may
## indicate that k is too low, especially if edf is close to k'.
##
##           k'      edf k-index p-value
## s(dateInt)    9.90e+01 1.01e+00   0.86   0.02 *
## s(monthFac)    1.20e+01 4.39e-05    NA    NA
## s(Avg_Temp)     9.00e+00 1.00e+00   1.06   0.80
## s(exchange_rate) 9.00e+00 1.01e+00   0.89   0.03 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```